



Project Five

ENEE 447: Operating Systems – Spring 2020

Assigned: Week of Mar 9; Due: Mar 27

Purpose

In this project, your main task is to implement non-blocking I/O for just two scenarios: reading and writing the DEV_WORD devices, which includes the LED (write-only) and the Console (read/write). You also need to implement the C-language scheduler, which is now more complex than in the previous project, as it handles a set of threads, not just two, and it maintains a TCB for each thread that is running. Much has been improved since the last project: you will work with a slightly more fleshed-out shell; you will invoke threads from the shell; and you can get a list of running processes. The thread creation and invocation facilities are more generalized and therefore support a more extensive system. Thus, you are being exposed to techniques that better exemplify the way things are normally done, and you will build the hardest parts (95% of the work has been done for you: the easy part that nonetheless would take quite a while to implement).

The Simplest Possible Operating System™

As said before, at its simplest, an OS is nothing more than a collection of vectors: it does nothing unless it is responding to interrupts. This is what we are building in Projects 3, 4, and 5. Project 3 implemented a system-call facility and provided an extremely rudimentary shell that runs in user mode — but still within the kernel proper, for now. Project 4 added to that the preemptive context switch and a rudimentary scheduler: the scheduler simply switches processes every time quantum, involved by a timer interrupt. Project 5 makes the I/O non-blocking so that the kernel can better respond to multiple sources of input “simultaneously” arriving from different sources.

Interrupt Handlers and Thread Context

Firstly, we have provided a simplified mechanism for saving and restoring state. In the *boot.s* code, you will find the following assembly-code routines:

```
save_lr_irq: .word 0
// courtesy of Prof Vince Weaver, U Maine
irq_handler:

    ldr    sp, tcb_address_runningthread
    stmia  sp, {r0-lr}^          @ Save all user registers r0-lr
                                   @ (the ^ means user registers)

    str    lr, [sp, #60]         @ store saved PC on stack
    str    lr, save_lr_irq      @ save the SVC lr
    mrs    lr, SPSR             @ load SPSR (assume ip not a swi arg)
    str    lr, [sp, #64]         @ store on stack
    ldr    lr, save_lr_irq      @ save the SVC lr
    @ Call the C version of the handler
    mov    sp, #SVCSTACK0
    bl     clear_timer_interrupt
    bl     periodic_timer
    bl     set_timer
    ldr    sp, tcb_address_runningthread
    ldr    r0, [sp, #64]         @ pop saved CPSR
    msr    SPSR_cxsf, r0        @ move it into place
    ldr    lr, [sp, #60]         @ restore address to return to

    @ Restore saved values. The ^ means to restore the userspace registers
    ldmia  sp, {r0-lr}^
    subs   pc, lr, #4           @ return from exception
```

```

save_r0_svc: .word 0
// courtesy of Prof Vince Weaver, U Maine
svc_handler:
    ldr    sp, tcb_address_runningthread
    stmia  sp, {r0-lr}^    @ Save all user registers r0-lr
    add    ip, lr, #4      @ (the ^ means user registers)
    str    ip, [sp, #60]   @ store saved PC on stack
    mrs    ip, SPSR        @ load SPSR (assume ip not a swi arg)
    str    ip, [sp, #64]   @ store on stack
    @ Call the C version of the handler
    mov    sp, #SVCSTACK0
    bl     trap_handler
    ldr    sp, tcb_address_runningthread
    ldr    r0, [sp, #64]   @ pop saved CPSR
    msr    SPSR_cxsf, r0  @ move it into place
    ldr    lr, [sp, #60]   @ restore address to return to

    @ Restore saved values. The ^ means to restore the userspace registers
    ldmia  sp, {r0-lr}^
    subs   pc, lr, #4      @ return from exception

```

These two routines are the only interrupt handlers, and the IRQ will neither interrupt itself (unless we set the timer too fast), nor will it interrupt the SVC mode. On a single core, the SVC mode will not interrupt itself. Therefore, none of this code will conflict with itself. What that means is that you won't have to worry about the *tcb_address_runningthread* value changing out from underneath you, and potential weirdness of that sort.

The code begins by performing the following save-register functions:

- The address of the TCB for the currently running thread is loaded into the *sp* register, which does not destroy the USER mode's copy of the *sp* register (see previous write-ups on the ARM register file).
- Registers r0–r14 are stored upwards starting at this address. These are the user registers, so the *sp* and *lr* registers are the user's copies.
- The return address is stored at the next address, which would correspond to the location for r15. This is because the return address is r15, as that *is* the Program Counter in the ARM32 architecture.
- Last, at the next location beyond that, we store the process's saved SPSR.

Therefore, one can think of the register set being saved as looking like the following:

```

REG_r0,
REG_r1,
REG_r2,
REG_r3,
REG_r4,
REG_r5,
REG_r6,
REG_r7,
REG_r8,
REG_r9,
REG_r10,
REG_r11,
REG_r12,
REG_sp,
REG_lr,
REG_pc,
REG_spsr,

```

That is exactly the data that is save and restored for a context switch. These values, in that order, are stored in the following structure, for which there is one for every process in the system:

```

struct tcb {
    LL_PTRS;
    char    name[NAMESIZE];
    long    threadid;
    long    stack;
    long    regs[17];    // 17th reg is the SPSR saved by context switch
} tcbs[ NUM_THREADS ];

```

Note that we have a statically-declared set of thread structures, *tcbs[**NUM_THREADS**]*. Thread 0 is the “NULL thread,” and the rest of the thread numbers can be assigned to anything.

Note that the *name*, *thread ID*, and *stack* are static values: they should not change. When a thread runs, all of its context is stored in the *regs* portion of the TCB. For instance, the *stack* member of the struct is the statically assigned starting point for the threads stack: user thread stacks begin at 0x20000, and each thread is given its own 4KB segment: thread 0 gets the first 4KB; thread 1 gets the next; etc. This will change once we have virtual memory, but just remember not to modify this value, and when a new thread starts up, use that as its initial stack pointer (for instance, you could put the value in *REG_sp*).

Back to the assembly code: Once the registers are saved, the code calls C-language routines to do the work. Because the handlers cannot be interrupted by themselves, we simply assign the same static starting stack location for each handler invocation. This works because we are running single-core for the moment.

After the C-language routines finish, we restore state from the thread control block:

- The address of the TCB for the currently running thread is loaded into the *sp* register.
- First, from the topmost location (corresponding to what would be r16 in the TCB), we restore the process’s saved SPSR.
- Next we retrieve the user process’s return address from the TCB location corresponding to r15. This is held temporarily in the *lr* link register, which is not the user link register.
- Lastly, registers r0–r14 are restored upwards starting at the bottom address. These are the user registers, so the *sp* and *lr* registers are the user’s copies and do not overwrite the “sp” or “lr” registers being used by the handler code.

Here’s the weirdness. The ARM32 architecture often used the following snippet to return from IRQ and FIQ exceptions:

```
subs    pc, lr, #4    @ return from exception
```

It is clearly a hack. It is inelegant. It is how IRQ and FIQ interrupts generally work. However—and here is where the weirdness comes in—it is NOT how ARM’s SVC trap handlers generally work. If you try to subtract 4 from the SVC handler’s link register on return from exception, you will get a stalled processor. But nonetheless, we would like to structure our kernel so that at both entry points into the kernel, a process is saved so that it can potentially be swapped out as a result of calling the interrupt.

What this requires is that the process must be saved and restored identically in both places, no matter which interrupt handler does the saving/restoring. Thus, what gets stored by the register-save protocol into the TCB is a number that will work with the “-4” on the return-from-exception at the end of both interrupt handlers (we do it in the SVC handler even though we don’t need to, just to ensure symmetry). This is expected for the IRQ, but it is not expected for the SVC trap handler. Thus, you see that, when the SVC handler stores the return value, it first adds 4 to it. Yes, this is inelegant. Yes, this is stupid. Welcome to interfacing with hardware. So what is stored in the TCB structure above is—and should be— 4 bytes larger than the PC destination where you really want to go. **Keep this in mind** when you are manually putting values into *REG_pc* or printing them out. I know that it is nonsensical, but it is a hack to overcome ARM32 weirdness.

The upside: you need not write any assembly code for this project, unless you want to.

General Kernel Structure

The idea is that on either a trap (SVC handler) or a periodic timer interrupt (IRQ handler), a process can get swapped out, so we save its context at the start and restore it at the end. This allows the kernel to put threads to sleep if they make sys calls that require blocking I/O, or to swap threads during the periodic timer interrupt. And, as mentioned before, these two interrupt handlers do not conflict with each other, so it should, in theory, all work out. Therefore, on a trap to the SVC handler, a process can get swapped out if it is a long-latency I/O operation (for this project, you will do this when encountering a `SYSCALL_RD_WORD` or `SYSCALL_WR_WORD`). On an invocation of the IRQ handler (due to a periodic timer interrupt), we will always swap processes, if there is another to swap to.

The *threads.c* module maintains three lists of processes:

- The *tfree* free list (contains TCB structures not being used)
- The *runq* (TCB structures for threads that *can* run)
- The *sleepq* (TCB structures for threads that cannot run and are waiting on I/O to finish)
- Pointers to the *active_thread* and *null_thread*

Note that a thread will be on one and only one of the queues at any given moment. In addition, the *tcbs[NUM_THREADS]* data block described earlier is indexed by the thread ID, so, for instance, if you want to find thread number *i* you can always do it this way:

```
struct tcb *tp = &tcbs[i];
```

You can manage these queues however you want ... for instance, should the NULL thread be on the *runq* or not? Should the currently executing thread (pointed to by the *active_thread* pointer) be on the *runq* or not? These are implementation details that are up to you and should reflect the way your scheduling algorithm is designed.

The *create_thread()* routine sets up a thread and moves its TCB structure onto the *runq*, where it is likely to begin executing (become the currently active thread), starting with the next periodic timer interrupt. It does not need to do anything more than this, because, for instance, in a given time quantum one could imagine multiple instances of *create_thread()* being called, and if each was allowed to think that it was selecting the next thread to run, they would all step on each other's toes.

The *scheduler()* function is what actually makes the selection of what thread to run, and make it the active thread. The function takes as an input one of three values:

- `THREAD_INIT` — indicates that there is no active thread to remove; we simply need to select someone from the *runq* and initialize the TCB pointer, as well as stack and start addresses that are going to be used by the *boot.s* assembly code (see where *core0* starts up the first user thread).
- `THREAD_RUN` — indicates that the currently active thread should be preemptively swapped out for another thread, provided that there is another to run. If there is no other thread on the *runq*, then the currently active thread should simply continue to execute.
- `THREAD_SLEEP` — indicates that the currently active thread needs to be put onto the *sleepq*, because it is waiting on long-latency I/O operations to finish. If the *runq* is empty, the NULL thread is run as a default; that is the only reason it exists.

Thus, if you want to put the current task to sleep, call *scheduler(THREAD_SLEEP)*; if you want to simply schedule the next available task on the *runq* but keep the currently executing thread active, call *scheduler(THREAD_RUN)*; etc.

The scheduler displaces whatever thread is currently active, moves it to the desired state (e.g., RUN or SLEEP), corresponding to moving it to one of the two queues, and then it selects a new thread to run and makes it active, which includes telling the interrupt handlers where to restore the thread's TCB state from.

Handling I/O Operations in a Non-Blocking Manner

As described in class, when your trap handler sees that the system call is SYSCALL_RD_WORD or SYSCALL_WR_WORD, you invoke the timeout queue instead of performing a blocking call to I/O routines. The mechanism for using the timeout queue has been changed slightly. First, the *create()* routine looks like this:

```
void create_timeoutq_event(int timeout, int repeat, int max, pfv_t function, namenum_t
data)
```

The *timeout* is the time until the event fires; the *repeat* value, if non-zero, indicates when after that it should re-fire. The *max* value indicates the maximum number of repeats, and the *function/data* pair acts like before.

Another difference, as you will see if you look in the *callout.c* module, is that the input to the user-specified timeout function is the data structure itself:

```
ep = (struct event *)LL_DETACH(timeoutq, ep);

ep->go(ep);

if (ep->repeat_interval > 0 && ep->max_repeats-- > 0) {
    ep->timeout = ep->repeat_interval; insert_event(ep);
} else {
    LL_PUSH(freelist, ep);
}
```

Your user-defined function is given read/write access to the data structure that determines how it will be inserted back onto the *timeoutq*, or whether it will be put back on the *timeoutq* at all. This allows your user-defined timeout function to determine for itself if it wants to continue repeating. If it wants to continue, it can manually set the struct's *repeat* and *max* values to appropriate non-zero values, and if it wants to terminate, it sets them to zero. This is what you will implement in the function *do_dev_word* in the *io.c* module.

Other facilities that this code will use are the new read and write *check()* functions in the device structure. You will notice that the device table has been redefined slightly, with the addition of these two new structure members:

```
struct dev {
    char devname[8];
    int devtype;
    pfv_t init;
    pfi_t read;
    pfi_t write;
    pfi_t rcheck;
    pfi_t wcheck;
};
```

Each device-table entry now contains *rcheck* and *wcheck* members, which are used to check and see if the device is ready to read or write. These functions have been provided for you, at least for the UART device (the main device for which you will be using them in this project). They are in the *uart.c* module:

```
int uart_rcvcheck ( void ) {
    if (GET32(AUX_MU_LSR_REG) & 0x01) return(1);
    if (GET32(AUX_MU_LSR_REG) & 0x01) return(1);
    if (GET32(AUX_MU_LSR_REG) & 0x01) return(1);
    if (GET32(AUX_MU_LSR_REG) & 0x01) return(1);
    if (GET32(AUX_MU_LSR_REG) & 0x01) return(1);
}
```

```

        return(0);
    }

    int uart_sendcheck ( void ) {
        if (GET32(AUX_MU_LSR_REG) & 0x20) return(1);
        if (GET32(AUX_MU_LSR_REG) & 0x20) return(1);
        if (GET32(AUX_MU_LSR_REG) & 0x20) return(1);
        if (GET32(AUX_MU_LSR_REG) & 0x20) return(1);
        if (GET32(AUX_MU_LSR_REG) & 0x20) return(1);

        return(0);
    }

```

Using repetitive checks is a hack to improve responsiveness. Recall that this replaces a *while()* loop—the fact there was a long-lasting series of checks in a *while()* loop suggests that sometimes it may take a while [pun intended] to get the UART’s attention. If we only checked once, and failed, then we might have to wait potentially 1/100 to 1/10 of a second to try again ... might as well give it a bunch of tries (kind of like replacing the *while()* loop with a short *for()* loop that can be retried periodically), because it increases the likelihood of success *without* actually becoming a *blocking* action.

Your code, which will implement the *do_dev_word()* function, will make use of these functions, testing to see if the device is ready to read or write — if the call to a *ready()* function indicates the device is ready, perform the function, but if not, reschedule the task to try again in the near future. Whenever the *do_dev_word()* function wakes up, it checks to see if the word-granularity I/O device is ready. If the device is ready, it performs the I/O read/write function and sets the event structure not to repeat; however, if the device is *not* ready, the function sets the repeat values in the event structure appropriately to try again at a specified point in the future. For WRITE system calls, the data to output to the device will be in the *data* value of the *io* structure. For READ system calls, you need to figure out how to get the value back to the user process, which is currently sleeping on the *sleepq*. Simply waking it up by putting it onto the *runq* is not enough—you also need to transfer the read data, a single word, back to the thread when it wakes up. Hint: use the sleeping thread’s TCB structure!

Example Interaction with the Shell

The following shows how the interaction with the shell might go, indicating what sort of things might come next in terms of thread invocation, etc.

```

[c0|00:02.023] ...
[c0|00:02.025] System is booting, kernel cpuid = 00000000
[c0|00:02.030] Kernel version: [p5-solution, Wed Mar 13 22:38:45 EDT 2019]
[c0|00:02.037] create_thread:
[c0|00:02.039] NULL thread 00000000
[c0|00:02.043] stack = 00021000
[c0|00:02.046] start = 00000044
[c0|00:02.049] tcb = 0000DBB8
[c0|00:02.052] create_thread:
[c0|00:02.054] shell 00000001
[c0|00:02.057] stack = 00022000
[c0|00:02.060] start = 00002294
[c0|00:02.063] tcb = 0000DC1C
[c0|00:02.066] ...
[c0|00:02.068] Init complete. Please hit any key to continue.
<hit any key>
Running shell.
Available commands:
RUN = 004E5552
PS = 00005350
TIME = 454D4954
LED = 0044454C
LOG = 00474F4C
EXIT = 54495845

```

```

Please enter a command.
> PS
CMD PS
[c0|00:09.695] PS: Active processes ...
[c0|00:09.698] Dumping TCB for thread 00000001
[c0|00:09.703] shell 00000001
[c0|00:09.705] stack 00022000
[c0|00:09.708] tcb @ 0000DC1C
[c0|00:09.711] r0 00000001
[c0|00:09.714] r1 0000000A
[c0|00:09.716] r2 00005350
[c0|00:09.719] r3 00005350
[c0|00:09.722] r4 00021FD4
[c0|00:09.725] r5 00000000
[c0|00:09.728] r6 00021FD4
[c0|00:09.731] r7 00000009
[c0|00:09.733] r8 54495845
[c0|00:09.736] r9 454D4954
[c0|00:09.739] r10 00002C84
[c0|00:09.742] r11 00000000
[c0|00:09.745] r12 00021FCD
[c0|00:09.747] sp 00021FA4
[c0|00:09.750] lr 000025D0
[c0|00:09.753] pc 00001A44
[c0|00:09.756] spsr 60000150

```

At this point, there are two threads that have been created: the NULL thread and the shell. However, when I do a PS, the NULL thread doesn't show up, because, in my implementation, it isn't kept on the *runq*. That is an unimportant detail; as mentioned above, you can do whatever is easiest or most logical for you.

At this point, we can start up *do_blinker* by invoking it from the shell. The inputs for RUN are a short name (1–3 characters, so that it fits into a 4-byte word with a NULL terminator) and a starting address. Looking at the *kernel7.list* file, we find that *do_blinker* is located at address 0x00002194. Thus, we have the following:

```

Please enter a command.
> RUN BLK 0X2194
CMD_RUN [BLK, 00002194]
[c0|00:54.005] SYSCALL_START THREAD name = 004B4C42
[c0|00:54.009] SYSCALL_START THREAD addr = 00002194
[c0|00:54.014] BLK 00002194
[c0|00:54.017] create_thread:
[c0|00:54.019] BLK 00000002
[c0|00:54.022] stack = 00023000
[c0|00:54.025] start = 00002194
[c0|00:54.028] tcb = 0000DC80
periodic blinking pattern begins
Please enter a command.
>

```

At this point, the *do_blinker()* routine starts running in user mode. It gets invoked the next timer interrupt after the "RUN" command is sent over as a system call, but you could just as easily have the trap handler put the shell onto the *runq* and make the blinker active immediately. All up to you. But it blinks regularly in the 1,2,3,4,1,2,3, ... pattern, all while the shell is allowing us to input commands and receive values. That is the hallmark of non-blocking I/O: things that expect to happen at regular intervals are not interrupted or stalled by long-running I/O operations.

At this point, if we do a PS, we should see both threads:

```

Please enter a command.

```

```

> PS

CMD PS
[c0|00:58.115] PS: Active processes ...
[c0|00:58.118] Dumping TCB for thread 00000001
[c0|00:58.122] shell 00000001
[c0|00:58.125] stack 00022000
[c0|00:58.128] tcb @ 0000DC1C
[c0|00:58.131] r0 00000001
[c0|00:58.133] r1 0000000A
[c0|00:58.136] r2 00005350
[c0|00:58.139] r3 00005350
[c0|00:58.142] r4 00021FD4
[c0|00:58.145] r5 00000000
[c0|00:58.147] r6 00021FD4
[c0|00:58.150] r7 00000009
[c0|00:58.153] r8 54495845
[c0|00:58.156] r9 454D4954
[c0|00:58.159] r10 00002C84
[c0|00:58.161] r11 00000000
[c0|00:58.164] r12 00021FCD
[c0|00:58.167] sp 00021FA4
[c0|00:58.170] lr 000025D0
[c0|00:58.173] pc 00001A44
[c0|00:58.175] spsr 60000150
[c0|00:58.178] Dumping TCB for thread 00000002
[c0|00:58.183] BLK 00000002
[c0|00:58.185] stack 00023000
[c0|00:58.188] tcb @ 0000DC80
[c0|00:58.191] r0 00000003
[c0|00:58.193] r1 00022FE0
[c0|00:58.196] r2 00000008
[c0|00:58.199] r3 00000000
[c0|00:58.202] r4 00000001
[c0|00:58.205] r5 000AAE60
[c0|00:58.208] r6 037EA464
[c0|00:58.210] r7 00000004
[c0|00:58.213] r8 00000000
[c0|00:58.216] r9 00000000
[c0|00:58.219] r10 00000000
[c0|00:58.222] r11 00000000
[c0|00:58.224] r12 00000000
[c0|00:58.227] sp 00022FDC
[c0|00:58.230] lr 000021F0
[c0|00:58.233] pc 000019E0
[c0|00:58.236] spsr 80000150

Please enter a command.
> TIME
CMD TIME = [00000000 047DD3DD]

Please enter a command.
>

```

Meanwhile, the blinking continues, with the correct timing and 1,2,3,4,1,2,3... pattern.

Among other things, this should make it relatively clear that, as far as the operating system goes, there is little difference between two different threads running in two different processes and two different threads running in the same process. Multithreading via spawning at the function level would be a trivial extension. Perhaps in an upcoming project...

Bottom Line

This truly represents the The Simplest Possible Operating System™ in that, once you have this working, you have a small-scale version of a Unix-like single-core operating system. It truly is a complete OS, the heart of any kernel. The main things that a “full” version of Unix or Linux would add to this are virtual memory and a file system. Don’t worry; we will get to those. In the meantime, this is basically the entire kernel. You now understand operating systems. Next we will go multicore. Buckle up.

Build It, Load It, Run It

Once you have it working, show us.